

Predictive Modeling

Lesson 4a

Linear Regression and Regression Trees

Kevin Zollicoffer

10/06/2013

Introduction

The RStudio project files and accompanying artifacts, including the tex file that created this PDF, are publicly available on GitHub

<https://github.com/zollie/PASS-PredictiveModeling-LinearRegressionAndRegressionTrees>

Data Setup

I took the Excel spreadsheet and saved it as a CSV for easy import into R

```
> bh <- read.csv("~/R/PASS/PredictiveModeling/LinearRegressionAndRegressionTrees/BostonHousing.csv")
> bh$X <- NULL
> bh$X.1 <- NULL
> bh$X.2 <- NULL
> bh$X.3 <- NULL
> bh$X.4 <- NULL
> head(bh)
```

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
1	0.00632	18	2.31	0	0.538	6.575	65.2	4.0900	1	296	15.3	396.90	4.98
2	0.02731	0	7.07	0	0.469	6.421	78.9	4.9671	2	242	17.8	396.90	9.14
3	0.02729	0	7.07	0	0.469	7.185	61.1	4.9671	2	242	17.8	392.83	4.03
4	0.03237	0	2.18	0	0.458	6.998	45.8	6.0622	3	222	18.7	394.63	2.94
5	0.06905	0	2.18	0	0.458	7.147	54.2	6.0622	3	222	18.7	396.90	5.33
6	0.02985	0	2.18	0	0.458	6.430	58.7	6.0622	3	222	18.7	394.12	5.21

MEDV

```
1 24.0
2 21.6
3 34.7
4 33.4
5 36.2
6 28.7
```

Partitioning

Next, the data is partitioned into 60% Train and 40% Test sets. I set the RNG seed for reproducibility

```
> set.seed(21275)
> n <- nrow(bh)
> a <- sort(sample(1:n, floor(n*.6)))
> bh.train <- bh[a,]
> bh.test <- bh[-a,]
```

Linear Regression

A Linear Regression model is fit to the train data. Then the model is used to make predictions using the test data while calculating the standard error for the Root Standard Mean Error

```
> mlr <- lm(MEDV ~ ., data=bh.train)
> summary(mlr)
```

Call:

```
lm(formula = MEDV ~ ., data = bh.train)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-10.4756	-2.8634	-0.4371	1.6646	22.1601

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	27.888067	7.848307	3.553	0.00048	***
CRIM	-0.106956	0.072117	-1.483	0.13971	
ZN	0.049067	0.021667	2.265	0.02467	*
INDUS	-0.022680	0.103597	-0.219	0.82695	
CHAS	4.518344	1.451186	3.114	0.00214	**
NOX	-13.912872	5.965888	-2.332	0.02075	*
RM	4.156809	0.634901	6.547	5.38e-10	***
AGE	-0.003258	0.022165	-0.147	0.88328	
DIS	-1.545232	0.331761	-4.658	6.02e-06	***
RAD	0.250664	0.100892	2.484	0.01384	*
TAX	-0.012236	0.005920	-2.067	0.04010	*
PTRATIO	-0.603644	0.219315	-2.752	0.00649	**
B	0.009321	0.004173	2.234	0.02668	*
LSTAT	-0.594316	0.081152	-7.324	6.78e-12	***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 4.735 on 189 degrees of freedom
Multiple R-squared: 0.7511, Adjusted R-squared: 0.734
F-statistic: 43.87 on 13 and 189 DF, p-value: < 2.2e-16

```
> mlr.pred <- predict(mlr, newdata=bh.test, se.fit=TRUE)
> head(mlr.pred)
```

```
$fit
      4      6      7     10     14     18     21     22
29.714160 25.957478 21.777154 17.219889 20.764395 17.644027 12.636732 18.441022
      28     31     35     39     42     45     46     47
15.237105 11.492917 14.056023 23.444027 28.588449 22.792016 21.834889 19.925409
      49     55     58     60     63     67     70     71
 7.017180 15.961282 32.702779 21.315283 24.472685 24.706993 20.821867 25.633714
      72     76     79     82     84     87     88     89
21.753825 23.886646 20.856452 27.503854 25.484312 22.159385 26.228971 31.443007
      94     96     97     98     99    100    102    109
29.237393 29.210982 24.781005 36.967139 36.412615 33.125572 27.055993 23.611240
      112    114    115    116    118    120    121    129
26.671856 20.117934 25.157003 19.902148 23.577127 20.415539 21.658533 19.696195
      130    132    133    136    137    138    142    143
14.337696 20.322249 21.063828 17.859595 16.371248 20.181758  2.956121 14.788836
      146    148    150    153    154    156    160    161
10.572343  6.529505 13.612512 21.228043 16.504114 21.414881 25.632117 33.647218
      167    169    171    172    174    180    181    185
36.988119 25.110348 20.859110 22.757070 28.922533 33.839084 35.645994 22.400581
      186    187    188    191    194    196    199    200
24.759640 37.131160 33.042259 30.436935 32.076026 40.932312 33.562971 30.584081
      201    208    209    213    219    223    224    225
31.203131 17.172166 25.039173 24.161332 25.420127 33.913340 29.746350 39.208190
      228    229    230    231    239    240    243    244
32.670880 36.316493 31.694445 23.992847 28.169229 28.051000 23.286563 27.178362
      246    248    249    250    251    254    255    259
12.438906 19.652636 21.329564 24.550954 24.690852 31.122316 23.563944 36.053773
      262    263    267    271    272    273    276    282
36.862921 41.065728 30.243691 22.277313 27.858715 29.059152 34.432398 33.911158
      283    285    286    287    292    294    295    296
42.048737 31.316309 26.682996 19.979341 35.689282 24.797868 23.147904 27.707513
      297    301    302    306    307    310    314    317
26.411145 29.987932 28.219071 31.280262 36.488668 23.781892 25.948645 17.118145
      320    321    322    328    330    331    332    334
21.382657 25.699561 25.679832 19.618868 24.293379 21.340947 19.349591 23.325918
      335    340    343    344    347    348    350    352
22.662304 22.123640 21.925966 28.091464 14.560947 25.581120 23.389614 20.868084
      357    361    364    366    368    370    373    376
21.281707 23.230933 21.983244 13.771885  9.625910 34.933418 27.839945 25.475565
```

378	383	386	389	390	393	394	395
19.661775	12.346955	6.414429	4.659947	13.260512	8.493454	19.941827	17.471469
399	401	402	403	405	409	410	411
4.966028	10.725885	17.217942	17.679711	6.095118	12.077657	19.108860	14.812751
414	417	421	425	427	431	432	435
10.458996	12.502894	19.470933	13.634894	15.602735	17.346477	17.874907	15.727103
436	438	440	446	452	454	455	456
12.819864	7.888171	12.228805	11.331696	19.127613	22.521070	14.959344	15.499756
457	462	465	470	472	473	478	479
12.098358	20.026672	20.059396	17.810058	22.161287	21.869056	9.939112	18.378290
481	482	484	485	487	490	491	493
22.887545	26.868539	20.404695	18.826205	18.863572	7.369337	2.297218	16.046951
498	500	501					
19.334566	18.556162	20.751059					

\$se.fit

4	6	7	10	14	18	21	22
0.9317122	0.9286232	0.9094579	1.3495501	0.9596900	0.8568180	1.0481473	0.9126928
28	31	35	39	42	45	46	47
0.9851723	1.0701011	1.1519041	0.9932167	1.1800970	0.7483313	0.8446413	0.9018019
49	55	58	60	63	67	70	71
1.7794641	1.5922848	1.5564002	0.9109190	1.1569029	1.2860336	0.7725690	1.1564554
72	76	79	82	84	87	88	89
1.0128518	0.8096550	0.8126226	0.7673139	0.6196457	0.7741901	0.7510598	0.9635369
94	96	97	98	99	100	102	109
1.3273137	0.8869292	0.9095414	1.1938209	1.2027777	0.9482638	0.9636533	1.0114876
112	114	115	116	118	120	121	129
0.7861166	0.8913017	0.8514661	0.8658377	0.8390837	0.8693890	2.0296093	1.2028232
130	132	133	136	137	138	142	143
1.1014328	1.2365316	1.2677682	1.1732646	1.1102194	1.2100763	1.5857398	1.8387206
146	148	150	153	154	156	160	161
1.6867581	1.6959211	1.4868280	1.8779506	1.5288533	2.0221573	1.7158873	1.7186939
167	169	171	172	174	180	181	185
1.5950180	1.2528310	1.2415745	1.2755666	0.9358624	0.9258309	1.1054596	1.1862111
186	187	188	191	194	196	199	200
0.9103022	1.0603503	1.2290780	1.0142139	1.0315543	1.3364976	1.3680617	1.4427962
201	208	209	213	219	223	224	225
1.4590041	0.7923440	1.5322894	1.5936021	1.6335361	1.5200233	0.7754441	1.1506324
228	229	230	231	239	240	243	244
0.8031847	1.3200924	1.1129999	0.7363773	0.7528767	0.6163152	0.8381933	0.9002992
246	248	249	250	251	254	255	259
1.2682497	1.3727987	0.9264054	0.9588302	0.9070551	1.7184301	1.2578433	1.3560648
262	263	267	271	272	273	276	282
1.3009348	1.4853447	1.3177693	0.8587378	1.0828205	0.7609909	0.9198071	0.8714906
283	285	286	287	292	294	295	296
1.6319925	1.2964835	0.8933280	1.3238859	1.6415385	1.2856638	1.1934003	1.2735060

```

      297      301      302      306      307      310      314      317
1.2809873 1.1569262 0.7263730 1.1260704 1.1987212 0.6510009 0.7153293 0.6890953
      320      321      322      328      330      331      332      334
0.5460082 0.6508959 0.6595567 0.7389261 1.5305962 1.4041473 0.8229418 1.1700935
      335      340      343      344      347      348      350      352
1.1470084 0.9609671 1.6945314 0.8575222 1.3107864 1.2464621 1.3070139 1.5901144
      357      361      364      366      368      370      373      376
1.5813572 1.3553720 1.5945469 2.3036478 1.9795235 1.7539784 1.6002766 1.3011771
      378      383      386      389      390      393      394      395
1.0419006 0.9365723 1.2273389 1.1725238 0.9367140 1.0314995 0.8582648 0.8736632
      399      401      402      403      405      409      410      411
2.1448516 1.3834186 0.9098533 0.8568119 2.2633278 1.1588585 1.0855765 3.1845183
      414      417      421      425      427      431      432      435
1.5001348 1.5206493 0.8170834 1.5613504 1.4655408 1.2358307 1.3200201 1.0956961
      436      438      440      446      452      454      455      456
1.1792701 1.4020738 0.9782476 1.3342367 0.9675572 1.2111130 1.4224180 1.3642674
      457      462      465      470      472      473      478      479
1.4867771 1.0244533 0.9380025 1.0017059 1.1757307 0.9986555 0.9628734 0.8341819
      481      482      484      485      487      490      491      493
1.0452997 1.1415572 1.2419500 1.1851493 0.9518138 1.8225458 1.9199579 1.8731892
      498      500      501
0.7372024 0.8533092 0.7142648

```

```
$df
```

```
[1] 189
```

```
$residual.scale
```

```
[1] 4.734828
```

```
> summary(mlr.pred$se.fit)
```

```

      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
0.5460   0.9097   1.1560   1.1870   1.3660   3.1850

```

```
> se <- mean(mlr.pred$se.fit)
```

```
> se
```

```
[1] 1.186619
```

```
> rse <- sqrt(se)
```

```
> rse
```

```
[1] 1.08932
```

Regression Trees

First a function is created to build Regression Trees with varying parameters.
Next, this function is used to build Regression Trees.

```

> f <- function(minspl, minbuck) {
+   require(rpart)
+   require(rpart.plot)
+
+   rtree <- rpart(MEDV ~., data=bh.train, method="anova", minsplit=minspl, minbucket=minbuck)
+
+   # pruned with min xerror
+   rtree.pruned <- prune(rtree, rtree$cptable[which.min(rtree$cptable[, "xerror"]), "CP"])
+ }
> rtree.A <- f(3, 1)
> rtree.B <- f(10, 3)
> rtree.C <- f(30, 1)
> printcp(rtree.A)

```

Regression tree:

```

rpart(formula = MEDV ~ ., data = bh.train, method = "anova",
      minsplit = minspl, minbucket = minbuck)

```

Variables actually used in tree construction:

```

[1] CRIM  DIS   LSTAT RM

```

Root node error: 25689/303 = 84.781

n= 303

	CP	nsplit	rel error	xerror	xstd
1	0.479933	0	1.00000	1.01169	0.109182
2	0.170368	1	0.52007	0.61426	0.073748
3	0.050214	2	0.34970	0.43368	0.059815
4	0.044850	3	0.29949	0.42787	0.063684
5	0.036171	4	0.25464	0.39895	0.063174
6	0.029695	5	0.21846	0.35276	0.062554
7	0.023446	6	0.18877	0.34427	0.062468

```

> rsq.rpart(rtree.A)

```

Regression tree:

```

rpart(formula = MEDV ~ ., data = bh.train, method = "anova",
      minsplit = minspl, minbucket = minbuck)

```

Variables actually used in tree construction:

```

[1] CRIM  DIS   LSTAT RM

```

Root node error: 25689/303 = 84.781

n= 303

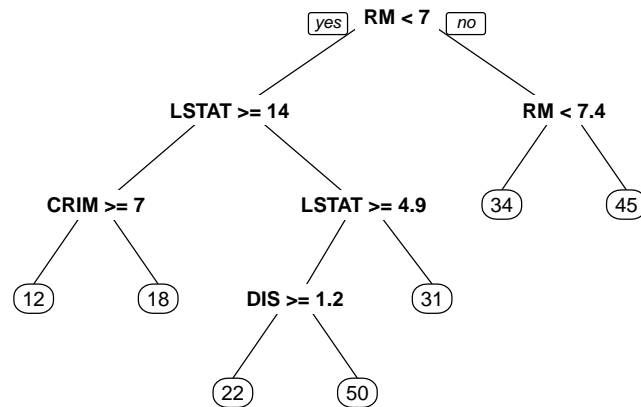


Figure 1: rtree.A

	CP	nsplit	rel error	xerror	xstd
1	0.479933	0	1.00000	1.01169	0.109182
2	0.170368	1	0.52007	0.61426	0.073748
3	0.050214	2	0.34970	0.43368	0.059815
4	0.044850	3	0.29949	0.42787	0.063684
5	0.036171	4	0.25464	0.39895	0.063174
6	0.029695	5	0.21846	0.35276	0.062554
7	0.023446	6	0.18877	0.34427	0.062468

```
> #plotcp(rtree.A)
```

```
> printcp(rtree.B)
```

Regression tree:

```
rpart(formula = MEDV ~ ., data = bh.train, method = "anova",
      minsplit = minspl, minbucket = minbuck)
```

Variables actually used in tree construction:

```
[1] CRIM  LSTAT RM
```

```
Root node error: 25689/303 = 84.781
```

```
n= 303
```

	CP	nsplit	rel error	xerror	xstd
1	0.479933	0	1.00000	1.00559	0.108379
2	0.170368	1	0.52007	0.59992	0.072855
3	0.050214	2	0.34970	0.42661	0.059573
4	0.044850	3	0.29949	0.42154	0.062043
5	0.036171	4	0.25464	0.38942	0.062886
6	0.019601	5	0.21846	0.34383	0.061656
7	0.016124	6	0.19886	0.32647	0.061155

```
> rsq.rpart(rtree.B)
```

```
Regression tree:
```

```
rpart(formula = MEDV ~ ., data = bh.train, method = "anova",  
      minsplit = minspl, minbucket = minbuck)
```

```
Variables actually used in tree construction:
```

```
[1] CRIM  LSTAT RM
```

```
Root node error: 25689/303 = 84.781
```

```
n= 303
```

	CP	nsplit	rel error	xerror	xstd
1	0.479933	0	1.00000	1.00559	0.108379
2	0.170368	1	0.52007	0.59992	0.072855
3	0.050214	2	0.34970	0.42661	0.059573
4	0.044850	3	0.29949	0.42154	0.062043
5	0.036171	4	0.25464	0.38942	0.062886
6	0.019601	5	0.21846	0.34383	0.061656
7	0.016124	6	0.19886	0.32647	0.061155

```
> #plotcp(rtree.B)
```

```
> printcp(rtree.C)
```

```
Regression tree:
```

```
rpart(formula = MEDV ~ ., data = bh.train, method = "anova",  
      minsplit = minspl, minbucket = minbuck)
```

```
Variables actually used in tree construction:
```

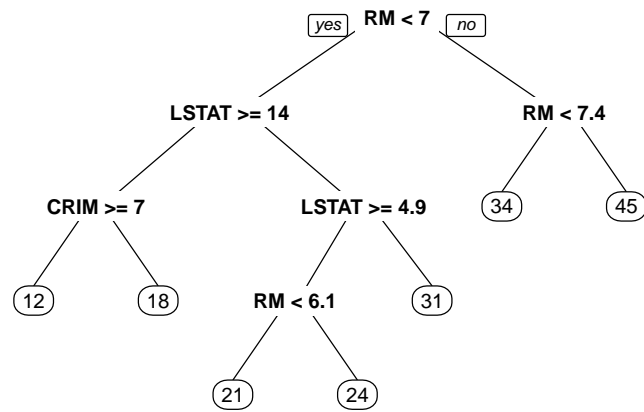



Figure 2: `rtree.B`

```
[1] CRIM DIS LSTAT RM
```

```
Root node error: 25689/303 = 84.781
```

```
n= 303
```

	CP	nsplit	rel error	xerror	xstd
1	0.479933	0	1.00000	1.00489	0.108849
2	0.170368	1	0.52007	0.58880	0.072471
3	0.050214	2	0.34970	0.45249	0.062622
4	0.044850	3	0.29949	0.42235	0.061164
5	0.036171	4	0.25464	0.39002	0.062550
6	0.029695	5	0.21846	0.36871	0.062673
7	0.015811	6	0.18877	0.35734	0.062362

```
> rsq.rpart(rtree.C)
```

```
Regression tree:
```

```
rpart(formula = MEDV ~ ., data = bh.train, method = "anova",  
      minsplit = minspl, minbucket = minbuck)
```

```
Variables actually used in tree construction:
```

```
[1] CRIM DIS LSTAT RM
```

```
Root node error: 25689/303 = 84.781
```

```
n= 303
```

	CP	nsplit	rel error	xerror	xstd
1	0.479933	0	1.00000	1.00489	0.108849
2	0.170368	1	0.52007	0.58880	0.072471
3	0.050214	2	0.34970	0.45249	0.062622
4	0.044850	3	0.29949	0.42235	0.061164
5	0.036171	4	0.25464	0.39002	0.062550
6	0.029695	5	0.21846	0.36871	0.062673
7	0.015811	6	0.18877	0.35734	0.062362

```
> #plotcp(rtree.C)
```

Lesson 4a Question and Answer

1

Report the root mean squared error on the validation data

```
> # root standard mean error from above
```

```
> rse
```

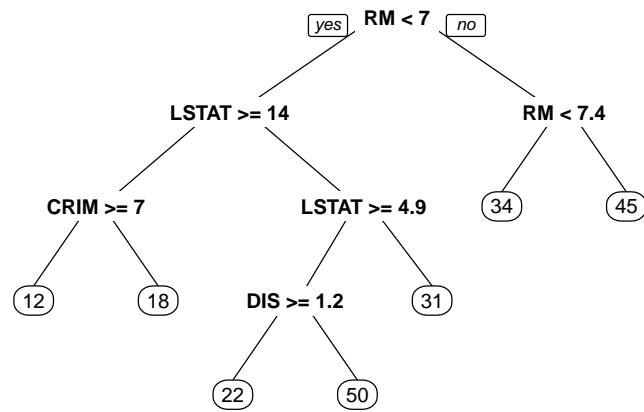


Figure 3: rtree.C

[1] 1.08932

2

Use the regression tree procedure in XLMiner to develop several models to predict the median value of houses in census tracts. Try multiple combinations of the tuning parameters

See Above.