

Regression Final

Kevin Zollicoffer

11/04/2013

Introduction

Regression final assignment using R

The complete source for this assignment is available on Github:

<https://github.com/zollie/PASS>

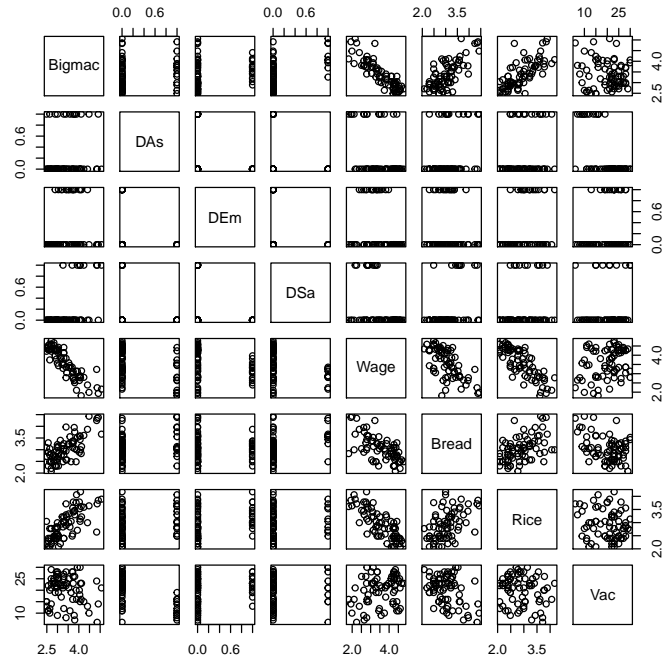
Problem 5.5

```
> ubs <- read.csv("~/R/PASS/Regression/RegressionFinal/ubs.csv")
> ubs$Region = factor(ubs$Region)
> head(ubs)
```

	City	Region	DAs	DEm	DSa	Bigmac	Wage	Bread	Rice	Vac
1	Amsterdam	NaWeOc	0	0	0	2.944	4.263	2.303	2.398	26
2	Athens	NaWeOc	0	0	0	3.401	3.829	2.565	3.296	22
3	Auckland	NaWeOc	0	0	0	2.944	3.786	2.944	2.565	21
4	Bangkok	Asia	1	0	0	3.807	2.653	3.761	3.296	7
5	Barcelona	NaWeOc	0	0	0	3.045	4.119	2.833	2.079	23
6	Beijing	Asia	1	0	0	3.784	2.625	3.951	3.434	9

a

```
> pairs(cbind(ubs$Bigmac,
+             ubs$DAs, ubs$DEm, ubs$DSa,
+             ubs$Wage, ubs$Bread,
+             ubs$Rice, ubs$Vac),
+       labels=cbind("Bigmac",
+                    "DAs", "DEm", "DSa",
+                    "Wage", "Bread",
+                    "Rice", "Vac"))
```



There seems to be clear linear relationships with Wage and Bigmac (negative), Bread and Bigmac, and Rice and Bigmac. The predictive power for Vac of Bigmac is less clear.

There seems to be some correlation between Wage and Bread, and Wage and Rice.

b

```
> model_b <- lm(Bigmac ~ DAs+DEm+DSa+Wage+Bread+Rice+Vac, data=ubs)
> stu_b <- rstudent(model_b)
> summary(model_b)
```

Call:

```
lm(formula = Bigmac ~ DAs + DEm + DSa + Wage + Bread + Rice +
    Vac, data = ubs)
```

Residuals:

Min	1Q	Median	3Q	Max
-0.44853	-0.16851	-0.03561	0.14513	0.63157

Coefficients:

Estimate	Std. Error	t value	Pr(> t)

```

(Intercept)  3.929449    0.631421    6.223 3.99e-08 ***
DAs          -0.207525    0.114411   -1.814  0.0743  .
DEm          -0.104153    0.094352   -1.104  0.2737
DSa           0.181246    0.118389    1.531  0.1306
Wage         -0.576032    0.084254   -6.837 3.37e-09 ***
Bread         0.309707    0.074294    4.169 9.25e-05 ***
Rice          0.102046    0.088928    1.148  0.2554
Vac           0.015415    0.005814    2.651  0.0101  *

```


Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.2331 on 65 degrees of freedom
Multiple R-squared: 0.8807, Adjusted R-squared: 0.8678
F-statistic: 68.54 on 7 and 65 DF, p-value: < 2.2e-16

```

> plot(ubs$Bread, stu_b)
> scatter.smooth(ubs$Bread, stu_b, span=.75)

```



The zero mean assumption can be tested with a loess line as above.


The constant variance assumption can be checked by ensuring the studentized

residuals are distributed evenly about the mean. That is, they have a mean of 0 and a standard deviation of 1.

The independence assumption can be checked with a quick glance at a residual scatter plot to ensure there are no non-random patterns.


The normality assumption can be checked with histograms and QQ plots as below.

```
> hist(stu_b)
```



To check the constant variance assumption we can use a histogram ...

```
> qqnorm(stu_b)
> qqline(stu_b)
```



c

```
> DasWage <- ubs$DAS*ubs$Wage
> DasBread <- ubs$DAS*ubs$Bread
> DasRice <- ubs$DAS*ubs$Rice
> DasVac <- ubs$DAS*ubs$Vac
> DemWage <- ubs$DEm*ubs$Wage
> DemBread <- ubs$DEm*ubs$Bread
> DemRice <- ubs$DEm*ubs$Rice
> DemVac <- ubs$DEm*ubs$Vac
> DsaWage <- ubs$DSa*ubs$Wage
> DsaBread <- ubs$DSa*ubs$Bread
> DsaRice <- ubs$DSa*ubs$Rice
> DsaVac <- ubs$DSa*ubs$Vac

> model_c <- lm(Bigmac ~ DAS+DEm+DSa+Wage+Bread+Rice
+               +Vac+DasWage+DasBread+DasRice+DasVac
+               +DemWage+DemBread+DemRice+DemVac+DsaWage
+               +DsaBread+DsaRice+DsaVac, data=ubs)
> summary(model_c)

Call:
lm(formula = Bigmac ~ DAS + DEm + DSa + Wage + Bread + Rice +
```

```
Vac + DasWage + DasBread + DasRice + DasVac + DemWage + DemBread +
DemRice + DemVac + DsaWage + DsaBread + DsaRice + DsaVac,
data = ubs)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.32656	-0.12469	-0.01427	0.10780	0.49710

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.683085	0.996793	3.695	0.000522	***
DAs	-0.738972	1.459567	-0.506	0.614749	
DEm	0.555744	1.383537	0.402	0.689532	
DSa	-0.984978	2.907659	-0.339	0.736134	
Wage	-0.478023	0.154255	-3.099	0.003106	**
Bread	0.082348	0.107430	0.767	0.446765	
Rice	0.221956	0.122558	1.811	0.075804	.
Vac	0.022065	0.007127	3.096	0.003131	**
DasWage	-0.103942	0.188635	-0.551	0.583933	
DasBread	0.567501	0.183594	3.091	0.003176	**
DasRice	-0.463639	0.202251	-2.292	0.025882	*
DasVac	0.054241	0.021818	2.486	0.016104	*
DemWage	-0.051375	0.214087	-0.240	0.811278	
DemBread	0.208266	0.175752	1.185	0.241303	
DemRice	-0.011267	0.203540	-0.055	0.956064	
DemVac	-0.045066	0.015387	-2.929	0.005005	**
DsaWage	-0.130008	0.460941	-0.282	0.779005	
DsaBread	0.690651	0.189595	3.643	0.000614	***
DsaRice	-0.160711	0.384538	-0.418	0.677684	
DsaVac	-0.007723	0.013443	-0.575	0.568045	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1903 on 53 degrees of freedom

Multiple R-squared: 0.9351, Adjusted R-squared: 0.9119

F-statistic: 40.21 on 19 and 53 DF, p-value: < 2.2e-16

The 3 interactions with largest p-value are ...

1. DemWage
2. DemRice
3. DsaWage

d

```
> model_d <- lm(Bigmac ~ DAs+DEm+DSa+Wage+Bread+Rice+Vac+DasWage+DasBread+DasRice+DasVac+DemBread+DemVac+DsaBread+DsaRice+DsaVac, data = ubs)
> summary(model_d)
```

Call:

```
lm(formula = Bigmac ~ DAs + DEm + DSa + Wage + Bread + Rice +
    Vac + DasWage + DasBread + DasRice + DasVac + DemBread +
    DemVac + DsaBread + DsaRice + DsaVac, data = ubs)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.32589	-0.12401	-0.01427	0.10882	0.49611

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)	
(Intercept)	3.873477	0.707313	5.476	1.06e-06	***
DAs	-0.929363	1.256451	-0.740	0.46259	
DEm	0.282447	0.565201	0.500	0.61922	
DSa	-1.752632	0.898657	-1.950	0.05615	.
Wage	-0.511492	0.100893	-5.070	4.66e-06	***
Bread	0.075781	0.102562	0.739	0.46306	
Rice	0.209768	0.092579	2.266	0.02735	*
Vac	0.022123	0.006832	3.238	0.00203	**
DasWage	-0.070473	0.146160	-0.482	0.63157	
DasBread	0.574068	0.177613	3.232	0.00206	**
DasRice	-0.451451	0.182008	-2.480	0.01616	*
DasVac	0.054183	0.021215	2.554	0.01340	*
DemBread	0.224488	0.158454	1.417	0.16210	
DemVac	-0.045118	0.013701	-3.293	0.00172	**
DsaBread	0.714585	0.166165	4.300	6.89e-05	***
DsaRice	-0.076633	0.179529	-0.427	0.67112	
DsaVac	-0.007269	0.012880	-0.564	0.57475	

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1854 on 56 degrees of freedom

Multiple R-squared: 0.935, Adjusted R-squared: 0.9164

F-statistic: 50.33 on 16 and 56 DF, p-value: < 2.2e-16

```
> anova(model_d, model_c)
```

Analysis of Variance Table

Model 1: Bigmac ~ DAs + DEm + DSa + Wage + Bread + Rice + Vac + DasWage +
DasBread + DasRice + DasVac + DemBread + DemVac + DsaBread +
DsaRice + DsaVac

```

Model 2: Bigmac ~ DAs + DEm + DSa + Wage + Bread + Rice + Vac + DasWage +
  DasBread + DasRice + DasVac + DemWage + DemBread + DemRice +
  DemVac + DsaWage + DsaBread + DsaRice + DsaVac
Res.Df    RSS Df Sum of Sq    F Pr(>F)
1      56 1.9243
2      53 1.9198  3 0.0045205 0.0416 0.9886

```

The p-value of .9886 > .05 therefore we cannot reject H_0 . We can remove the 3 interactions with the largest p-value as statistically $DemWage = DemRice = DsaWage = 0$

e

```

> model_e <- lm(Bigmac ~ DAs+DEm+DSa+Wage+Bread
+               +Rice+Vac+DasBread+DasRice+DasVac
+               +DemBread+DemVac+DsaBread, data=ubs)
> summary(model_e)

```

Call:

```

lm(formula = Bigmac ~ DAs + DEm + DSa + Wage + Bread + Rice +
  Vac + DasBread + DasRice + DasVac + DemBread + DemVac + DsaBread,
  data = ubs)

```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.31559	-0.11941	-0.00559	0.10736	0.49558

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.110961	0.553962	7.421	5.23e-10 ***
DAs	-1.512860	0.632105	-2.393	0.019894 *
DEm	0.212223	0.542260	0.391	0.696935
DSa	-2.219990	0.508840	-4.363	5.23e-05 ***
Wage	-0.541401	0.071295	-7.594	2.67e-10 ***
Bread	0.065805	0.098576	0.668	0.507019
Rice	0.189425	0.073788	2.567	0.012809 *
Vac	0.020716	0.005185	3.995	0.000182 ***
DasBread	0.608153	0.163987	3.709	0.000463 ***
DasRice	-0.394569	0.137903	-2.861	0.005830 **
DasVac	0.057814	0.019994	2.892	0.005358 **
DemBread	0.226347	0.154589	1.464	0.148451
DemVac	-0.042906	0.012967	-3.309	0.001600 **
DsaBread	0.729123	0.153530	4.749	1.35e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1816 on 59 degrees of freedom
Multiple R-squared: 0.9342, Adjusted R-squared: 0.9197
F-statistic: 64.47 on 13 and 59 DF, p-value: < 2.2e-16

```
> anova(model_e, model_d)
```

Analysis of Variance Table

Model 1: Bigmac ~ DAs + DEm + DSa + Wage + Bread + Rice + Vac + DasBread +
DasRice + DasVac + DemBread + DemVac + DsaBread
Model 2: Bigmac ~ DAs + DEm + DSa + Wage + Bread + Rice + Vac + DasWage +
DasBread + DasRice + DasVac + DemBread + DemVac + DsaBread +
DsaRice + DsaVac

	Res.Df	RSS	Df	Sum of Sq	F	Pr(>F)
1	59	1.9464				
2	56	1.9243	3	0.022104	0.2144	0.886

The p-value of .886 > .05 therefore we cannot reject H_0 . We can remove the 3 interactions considered here.

f

```
> model_f <- lm(Bigmac ~ DAs+DEm+DSa+Wage+Bread+Rice+
+ Vac+DasBread+DasRice+DasVac+DemBread+DemVac+DsaBread, data=ubs)
> summary(model_f)
```

Call:

```
lm(formula = Bigmac ~ DAs + DEm + DSa + Wage + Bread + Rice +
    Vac + DasBread + DasRice + DasVac + DemBread + DemVac + DsaBread,
    data = ubs)
```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.31559	-0.11941	-0.00559	0.10736	0.49558

Coefficients:

	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	4.110961	0.553962	7.421	5.23e-10 ***
DAs	-1.512860	0.632105	-2.393	0.019894 *
DEm	0.212223	0.542260	0.391	0.696935
DSa	-2.219990	0.508840	-4.363	5.23e-05 ***
Wage	-0.541401	0.071295	-7.594	2.67e-10 ***
Bread	0.065805	0.098576	0.668	0.507019
Rice	0.189425	0.073788	2.567	0.012809 *
Vac	0.020716	0.005185	3.995	0.000182 ***
DasBread	0.608153	0.163987	3.709	0.000463 ***
DasRice	-0.394569	0.137903	-2.861	0.005830 **

```

DasVac      0.057814    0.019994    2.892 0.005358 **
DemBread    0.226347    0.154589    1.464 0.148451
DemVac      -0.042906    0.012967   -3.309 0.001600 **
DsaBread     0.729123    0.153530    4.749 1.35e-05 ***

```

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1816 on 59 degrees of freedom
Multiple R-squared: 0.9342, Adjusted R-squared: 0.9197
F-statistic: 64.47 on 13 and 59 DF, p-value: < 2.2e-16

$H_0 : DemBread = 0$ $H_a : DemBread \neq 0$

.1484 > .05 therefore we do not reject H_0

g

```

> model_g <- lm(Bigmac ~ DAs+DEm+DSa+Wage+Bread+Rice+
+ Vac+DasBread+DasRice+DasVac+DemVac+DsaBread, data=ubs)
> summary(model_g)

```

Call:

```

lm(formula = Bigmac ~ DAs + DEm + DSa + Wage + Bread + Rice +
    Vac + DasBread + DasRice + DasVac + DemVac + DsaBread, data = ubs)

```

Residuals:

	Min	1Q	Median	3Q	Max
	-0.27016	-0.12943	0.00022	0.10108	0.51102

Coefficients:


	Estimate	Std. Error	t value	Pr(> t)
(Intercept)	3.861757	0.532170	7.257	9.12e-10 ***
DAs	-1.207924	0.602473	-2.005	0.049487 *
DEm	0.876811	0.299504	2.928	0.004821 **
DSa	-1.991641	0.488946	-4.073	0.000138 ***
Wage	-0.547934	0.071831	-7.628	2.11e-10 ***
Bread	0.154348	0.078586	1.964	0.054162 .
Rice	0.195452	0.074372	2.628	0.010889 *
Vac	0.021595	0.005199	4.153	0.000105 ***
DasBread	0.515727	0.152786	3.375	0.001297 **
DasRice	-0.406481	0.138969	-2.925	0.004856 **
DasVac	0.056577	0.020165	2.806	0.006759 **
DemVac	-0.043681	0.013079	-3.340	0.001446 **
DsaBread	0.641239	0.142652	4.495	3.24e-05 ***

Signif. codes: 0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1

Residual standard error: 0.1834 on 60 degrees of freedom
 Multiple R-squared: 0.9318, Adjusted R-squared: 0.9182
 F-statistic: 68.36 on 12 and 60 DF, p-value: < 2.2e-16

```
> stu_g <- rstudent(model_g)
> plot(ubs$Bread, stu_g)

> scatter.smooth(ubs$Bread, stu_g, span=.75)
```



The zero mean regression assumption now appears met.

h

$$\hat{Bigmac} = 3.861757 - 1.2709D_{As} + .8768D_{Em} - 1.9916D_{Sa} - .5479Wage + .15348Bread + .195452Rice + .021595Vac + .515727DasBread - .406481DasRice + .056577DasVac - .043681DemVac + .641239DsaBread$$

i

Yes, the economist is correct. The coefficient for Wage is negative meaning as Wage increases \hat{Bigmac} decreases all else held constant.

j


```
> mWage <- mean(ubs$Wage)
> mBread <- mean(ubs$Bread)
> mRice <- mean(ubs$Rice)
> mVac <- mean(ubs$Vac)
```

$$\hat{Bigmac} = 3.861757 - 1.2709D_{As} + .8768D_{Em} - 1.99.16D_{Sa} - mWage * Wage + .15348 * Bread + mRice * Rice + mVac * Vac + .15348 * DasBread - mRice * DasRice + mVac * DasVac - mVac * DemVac + .15348 * DsaBread$$

```
> DasBread0 <- DasBread[ubs$DAs==0]
> DasBread1 <- DasBread[ubs$DAs==1]
> DsaBread0 <- DsaBread[ubs$DSa==0]
> DsaBread1 <- DsaBread[ubs$DSa==1]
> DasRice0 <- DasRice[ubs$DAs==0]
> DasRice1 <- DasRice[ubs$DAs==1]
```

***I am having trouble producing the predictor effect plots below*

```
> plot(ubs$Bread, DasBread, type="n")
> lines(sort(ubs$Bread[ubs$DasBread==0]), DasBread0[order(ubs$Bread[ubs$DasBread==0])])
> lines(sort(ubs$Bread[ubs$DasBread==1]), DasBread1[order(ubs$Bread[ubs$DasBread==1])])
>
```



k

$$\hat{Bigmac} = 3.861757 - 1.2709D_{As} + .8768D_{Em} - 1.99.16D_{Sa} - mWage * Wage + mBread * Bread + .195452 * Rice + mVac * Vac + mBread * DasBread - .195452 * DasRice + mVac * DasVac - mVac * DemVac + mBread * DsaBread$$

i

$$\hat{Bigmac} = 3.861757 - 1.2709D_{As} + .8768D_{Em} - 1.99.16D_{Sa} - mWage * Wage + mBread * Bread + mRice * Rice + .021595 * Vac + mBread * DasBread - mRice * DasRice + .021595 * DasVac - .021595 * DemVac + mBread * DsaBread$$